

Lab 07 — Build a /daily-brief Skill

Topic 02 · Tools and Skills

Meridian Capital Partners Pte Ltd — all data in this lab is synthetic and for training only.

Goal

Labs 4 and 5 showed what the connectors can do, and Lab 6 showed a built-in skill doing the work for you. Now you write one yourself. The reliable way to build a skill is the Do / Verify / Save pattern: do the task with Claude once, refine until the output is right, then use `/skill-creator` to save it. You have already done the doing — this lab turns it into a command, then updates it. Finally you schedule it, so the briefing is waiting as a draft before you sit down rather than something you have to remember to ask for.

You'll build

A daily-brief Agent Skill you can run by name each morning, updated once and scheduled to run itself at 9am.

Tools: Claude Cowork, Agent Skills, `/skill-creator`, Google Drive connector, Gmail connector, Scheduled tasks

The prompt

This lab has **three blocks**, sent at different times. Blocks 1 and 2 are typed into the chat composer — block 1 does the work, and block 2 saves it once you are happy with the output. Block 3 schedules it. All three are also in `PROMPT.md`.

Block 1 — DO

Send this first, on its own:

PROMPT — copy into Claude

Run my morning finance briefing for Meridian Capital Partners, once, so we can get the format right together:

- What changed in the Meridian Finance folder on Google Drive since yesterday?
- Which desks are behind their quarterly target right now?
- What finance emails do I need to respond to today?

If a section has nothing in it, say so plainly — "No new files since yesterday" — rather than padding it out. End with one line on what most needs my attention, or "Nothing urgent — have a good day" if there is nothing. All money in Singapore dollars. Draft only — never send email.

Now stop and read the output. Ask for changes until the format is right — reorder a section, tighten a heading, fix the fallback wording. Whatever you settle on here is exactly what the skill will reproduce every morning. Do not continue until you are happy with it.

Block 2 — SAVE

Only once the output is right, type `/skill-creator` into the composer. It appears in the slash-command list as you type — select it, then add your description after it:

PROMPT — copy into Claude

```
/skill-creator Save the morning finance briefing routine we just ran as a skill called daily-brief. Keep the three sections in the order we settled on, the plain fallback wording for empty sections, the single closing attention line, all money in Singapore dollars, and draft-only – never send email.
```

Note: `/skill-creator` is a slash command you type, the same way you will later run `/daily-brief`. It is not something Claude invokes for you from inside a longer prompt.

Block 3 — SCHEDULE

This one does not go in the chat. Open the **Create scheduled task** dialog and paste it into the **Instructions** field:

PROMPT — copy into Claude

```
Run my daily-brief skill for Meridian Capital Partners.
```

```
Leave the briefing as a draft email to me in Gmail, with the subject "Meridian daily brief" and today's date. Do not send it.
```

```
If a section has nothing in it, keep the plain fallback wording rather than padding it out.
```

Notice it **names the skill** rather than restating the routine — that is what saving the skill bought you. Notice too that it repeats **Do not send it** even though the skill is already draft-only. A scheduled task is the one place you state a safety rule twice: it runs unattended with Permissions set to **Automatically approve**, so both the instruction and the skill carry the constraint.

Note: Do it once, check the output, then save it — that order is what makes the skill reliable.

Steps

1. Make sure the Drive and Gmail connectors from Labs 4 and 5 are still enabled.

The skill drives those connectors, so it can only work if they are on.

2. DO — send block 1 on its own and let Claude run the routine once.

You are not writing instructions yet; you are doing the task together so Claude sees what good looks like.

3. VERIFY — read the briefing and ask for changes until the format is right.

Reorder a section, tighten the wording, fix a heading. Whatever you settle on here is what the skill will reproduce.

4. Check the fallback wording on any empty section.

The fallbacks are what stop the routine inventing filler on a quiet morning — get them right before you save.

5. Confirm it drafted and did not send anything.

A routine you will run unattended must not be able to send mail on its own.

6. SAVE — type `/skill-creator` in the composer, then send block 2.

You type the slash command yourself and pick it from the list; it then writes the SKILL.md for you, from the run you just verified.

7. Read the SKILL.md it produces, especially the description line.

The description is how Claude decides whether to load the skill for a future task — make it specific.

8. Run the skill by typing `/daily-brief`, then run it again.

The same structure both times is the payoff: one command instead of re-describing the routine every morning.

9. UPDATE — ask for one change, such as putting the emails section first.

A skill is not written once. You keep refining it as you learn what you actually want each morning.

10. Save the change under the same skill name and confirm the replace prompt.

Claude warns that a skill with this name already exists and that replacing it cannot be undone — read that prompt before you accept it.

11. Run `/daily-brief` once more to confirm the change took effect.

Always re-run after an update; that is how you know the new version is the one being loaded.

12. Open Scheduled tasks from the sidebar, choose New task, then Set up manually.

The other route, Create with Claude, fills the task in for you by asking questions; setting it up manually shows you every field it will set.

13. Name the task "Daily briefing".

The skill captures WHAT to do; the schedule captures WHEN. Two different kinds of reuse, and this is the second one.

14. Paste Block 3 into the Instructions field.

It names the skill you just built rather than re-describing the routine — that is what saving it bought you. It also repeats "Do not send it", because a scheduled task is the one place you state a safety rule twice.

15. Point "Work in a project or folder" at the folder you have been working in.

The task needs the same working folder and connectors the skill expects.

16. Set Frequency to run daily at 9:00am.

Before you sit down, so the briefing is waiting rather than something you go and fetch.

17. Read the Permissions setting and understand what "Automatically approve" means.

The task will use your connectors unattended, without pausing to ask. This is exactly why the skill drafts and never sends — an unattended task that could send mail is how an unreviewed email reaches a real client.

18. Decide whether to switch on "Require this computer".

On, it only runs while your computer is awake but can reach your local folders; off, it runs without them. Choose to match where the task's data actually lives.

19. Save the task, then run it once manually to confirm it works.

Never leave a scheduled task untested — the first time it runs should not be the morning you are relying on it.

20. Check Gmail for a draft titled "Meridian daily brief" and confirm nothing was sent.

A draft in your Drafts folder and an empty Sent folder is the proof the constraint held.

Check your work

Running `/daily-brief` produces a briefing with all three sections — Drive changes, desks behind target, emails needing a reply — using the plain fallback wording where a section is empty, ending with a single attention line, with all money in SGD and no email sent. After the update, the revised section order appears on the next run — and Claude warned you before replacing the existing skill. A scheduled task named "Daily briefing" exists, set to run daily at 9:00am, and running it manually leaves a Gmail draft titled "Meridian daily brief" with today's date, while your Sent folder stays empty.

Before you move on

Use only the supplied mock data and your own test accounts. Never put real client, personal or confidential data into a training exercise, and send lab emails to yourself rather than to a real recipient.